

Operating Systems Design

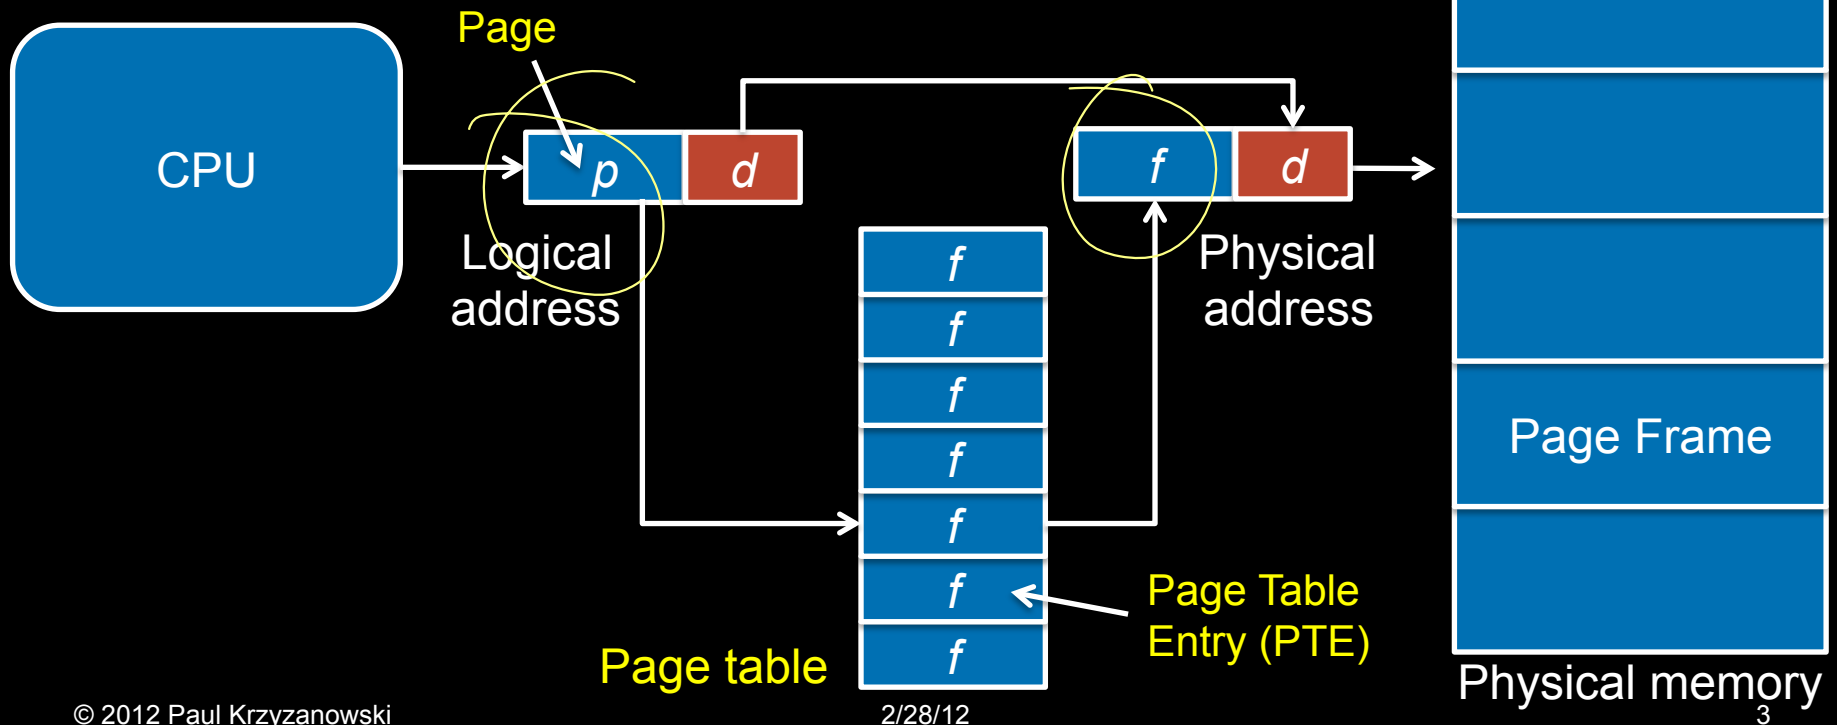
10. Memory Management: Paging



Paul Krzyzanowski
pxk@cs.rutgers.edu

Page translation



$$f = \text{page_table}[p]$$






Return oriented
prog.


"Buffer overflow
attack"

Virtual memory makes memory sharing easy




Sharing is by page granularity.

Virtual memory makes memory sharing easy



Copy on write

- Share until a page gets modified
- Example: fork()
 - Set all pages to read-only
 - Trap on write
 - If legitimate write
 - Allocate a new page and copy contents



Demand Paging

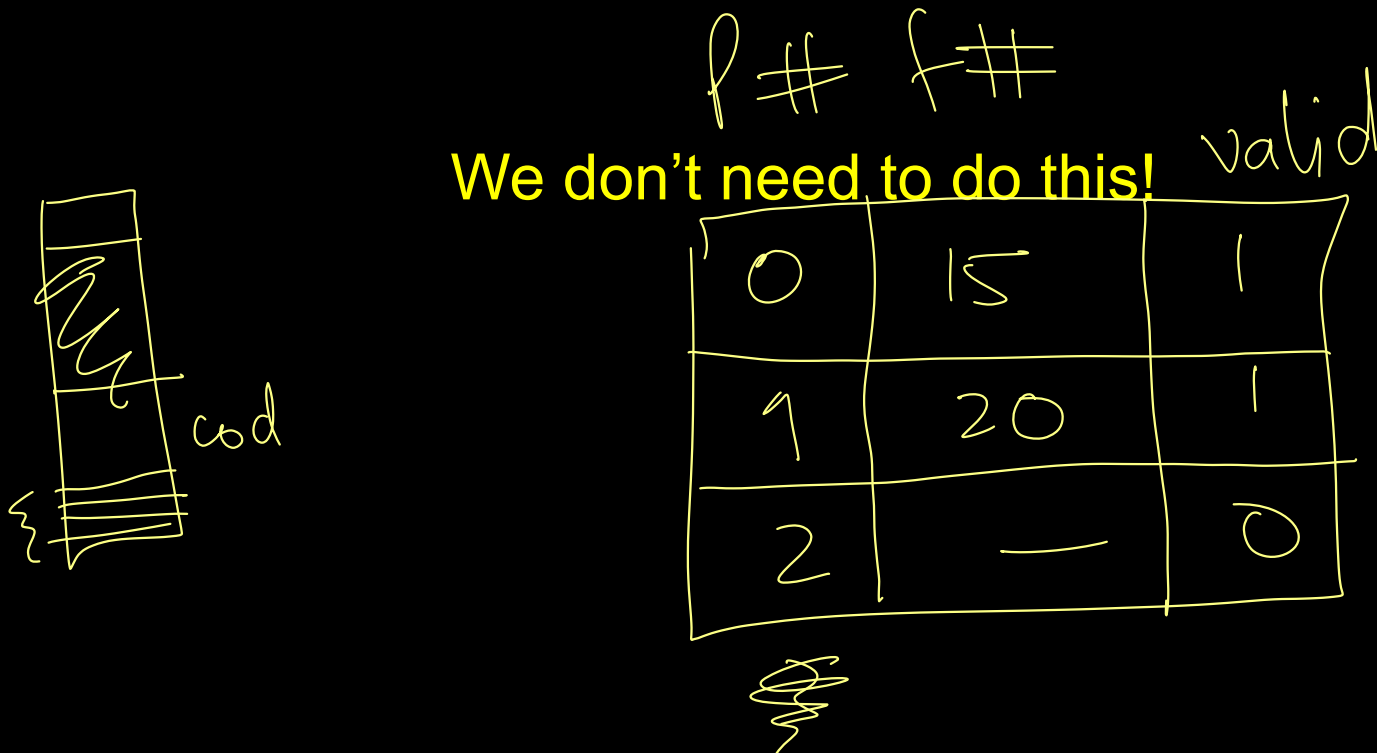


Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it

Executing a program

- Allocate memory + stack and load the entire program from memory (including linked libraries)
- Then execute it




Demand Paging

- Load pages into memory only as needed
 - On first access
 - Pages that are never used never get loaded
- Use valid/invalid bit in page table entry
 - Valid: the page is in memory (“valid” mapping)
 - Invalid: out of bounds access **or** page is not in memory
 - Have to check the process’ memory map in the PCB to find out
- Invalid memory access generates a ***page fault***

Demand Paging: At Process Start

- Open executable file
- Set up memory map (stack & text/data/bss)
 - But don't load anything!
- Load first page & allocate initial stack page
- Run it!



Page Replacement

- A process can run without having all of its memory allocated
 - It's allocated on demand
- If the
 $\{\text{address space used by all processes} + \text{OS}\} \leq \text{physical memory}$
then we're ok
- Otherwise:
 - Make room: discard or store a page onto the disk
 - If the page came from a file & was not modified
 - Discard ... we can always get it
 - If the page is dirty, it must be saved in a **swap file** / *partition / area /*
 - Swap file: a file (or disk partition) that holds excess pages

Virtual Memory

space.

Cost

- Handle page fault exception: ~ 400 usec
- Disk seek & read: ~ 10 msec
- Memory access: ~ 100 ns
- Page fault degrades performance by around 100,000!!
- Avoid page faults!
 - If we want < 10% degradation of performance, we can have just one page fault per 1,000,000 memory accesses

Page replacement

- We need a good replacement policy for good performance

FIFO Replacement

- First In, First Out
- Good
 - May get rid of initialization code or other code that's no longer used
- Bad
 - May get rid of a page holding frequently used global variables

Least Recently Used (LRU)


- Timestamp a page when it is accessed
- When we need to remove a page, search for the one with the oldest timestamp
- Nice algorithm but...
 - Timestamping is a pain – we can't do it with the MMU!

Not Frequently Used Replacement

- Approximate LRU
- Each PTE has a reference bit
- Keep a counter for each page frame
- At each clock interrupt:
 - Add the reference bit of each frame to its counter
 - Clear reference bit
- To evict a page, choose the frame with the lowest counter
- Problem
 - No sense of time: a page that was used a lot a long time ago may still have a high count
 - Updating counters is expensive

Thrashing

- Locality
 - Process migrates from one locality (working set) to another
- Thrashing
 - Occurs when sum of all working sets $>$ total memory



The End